



Introduction to TensorFlow Lite

Madeira International Workshop in Machine Learning


Diogo Freitas

July 30, 2021




SEEHEALTH

Where are we now?



Now we want to deploy our model! ☺ TensorFlow Lite

Where are we now?



Now we want to deploy our model! 🍷 TensorFlow Lite

Wait...



Why TensorFlow Lite?

- TensorFlow Lite enables **on-device machine learning** (e.g., Android and iOS devices, and microcontrollers).
- TensorFlow models can be used **without an internet connection**.
- Enables to create **lightweight models** with **low latency** (i.e., with high performance).
- Includes techniques for **hardware acceleration and model optimization**.
- Low power consumption.

Why TensorFlow Lite?

- TensorFlow Lite enables **on-device machine learning** (e.g., Android and iOS devices, and microcontrollers).
- TensorFlow models can be used **without an internet connection**.
- Enables to create **lightweight models** with **low latency** (i.e., with high performance).
- Includes techniques for **hardware acceleration and model optimization**.
- Low power consumption.

Why TensorFlow Lite?

- TensorFlow Lite enables **on-device machine learning** (e.g., Android and iOS devices, and microcontrollers).
- TensorFlow models can be used **without an internet connection**.
- Enables to create **lightweight models** with **low latency** (i.e., with high performance).
- Includes techniques for **hardware acceleration and model optimization**.
- Low power consumption.

Why TensorFlow Lite?

- TensorFlow Lite enables **on-device machine learning** (e.g., Android and iOS devices, and microcontrollers).
- TensorFlow models can be used **without an internet connection**.
- Enables to create **lightweight models** with **low latency** (i.e., with high performance).
- Includes techniques for **hardware acceleration and model optimization**.
- Low power consumption.

Why TensorFlow Lite?

- TensorFlow Lite enables **on-device machine learning** (e.g., Android and iOS devices, and microcontrollers).
- TensorFlow models can be used **without an internet connection**.
- Enables to create **lightweight models** with **low latency** (i.e., with high performance).
- Includes techniques for **hardware acceleration and model optimization**.
- Low power consumption.

1

Build a
model

2

Export and
convert

3

Verify

4

Deploy

1

Build a
model

2

Export and
convert

3

Verify

4

Deploy

Build a model

We already know how to build a model, but let us review it.

1. Create some dummy data (for demonstration purposes only):

```
1 X = np.array([1, 2, 3, 4, 5, 6])  
2 y = np.array([3, 5, 7, 9, 11, 13])
```

2. Create the model with a single hidden unit:

```
1 model = tf.keras.models.Sequential(  
2     [tf.keras.layers.InputLayer(input_shape=(1,)),  
3     tf.keras.layers.Dense(units=1)]  
4 )
```

Build a model

We already know how to build a model, but let us review it.

1. Create some dummy data (for demonstration purposes only):

```
1 X = np.array([1, 2, 3, 4, 5, 6])  
2 y = np.array([3, 5, 7, 9, 11, 13])
```

2. Create the model with a single hidden unit:

```
1 model = tf.keras.models.Sequential(  
2     [tf.keras.layers.InputLayer(input_shape=(1,)),  
3     tf.keras.layers.Dense(units=1)]  
4 )
```

3. Compile the model:

```
1 model.compile(optimizer='sgd', loss='mean_squared_error')
```

4. Train the model:

```
1 model.fit(X, y, epochs=10)
```

3. Compile the model:

```
1 model.compile(optimizer='sgd', loss='mean_squared_error')
```

4. Train the model:

```
1 model.fit(X, y, epochs=10)
```

1

Build a
model

2

Export and
convert

3

Verify

4

Deploy

You can generate a TensorFlow Lite model by:

- using an **existing TensorFlow** Lite model.
- **creating a TensorFlow Lite** model.
- **converting a TensorFlow model** into a TensorFlow Lite model.

File model extension:

Files generated by TensorFlow Lite are identified by the *.tflite* file extension.

You can generate a TensorFlow Lite model by:

- using an **existing TensorFlow** Lite model.
- **creating a TensorFlow Lite** model.
- **converting a TensorFlow model** into a TensorFlow Lite model.

File model extension:

Files generated by TensorFlow Lite are identified by the *.tflite* file extension.

You can generate a TensorFlow Lite model by:

- using an **existing TensorFlow** Lite model.
- **creating a TensorFlow Lite** model.
- **converting a TensorFlow model** into a TensorFlow Lite model.

File model extension:

Files generated by TensorFlow Lite are identified by the *.tflite* file extension.

You can generate a TensorFlow Lite model by:

- using an **existing TensorFlow** Lite model.
- **creating a TensorFlow Lite** model.
- **converting a TensorFlow model** into a TensorFlow Lite model.

File model extension:

Files generated by TensorFlow Lite are identified by the ***.tflite*** file extension.

Components of TensorFlow Lite

Converter

- Takes a TensorFlow model and generates a TensorFlow Lite model.
- Enables to export the model with various optimizations options and various platforms.

Interpreter

- Enables to run inference (i.e., given the inputs, compute the model's output) on a client device.
- Enables to run the model in various platforms (e.g., Linux and iOS).
- Provides hardware-accelerated APIs.

Components of TensorFlow Lite

Converter

- Takes a TensorFlow model and **generates a TensorFlow Lite model**.
- Enables to export the model with **various optimizations options** and various platforms.

Interpreter

- Enables to **run inference** (i.e., given the inputs, compute the model's output) on a client device.
- Enables to run the model **in various platforms** (e.g., Linux and iOS).
- Provides **hardware-accelerated APIs**.

1. Export the *SavedModel*¹:

```
1 export_dir = '/tmp/saved_model/1'  
2 tf.saved_model.save(model, export_dir=export_dir)
```

2. Convert the model:

```
1 converter = tf.lite.TFLiteConverter.from_saved_model(  
    export_dir)  
2 tflite_model = converter.convert()
```

3. Save the model:

```
1 tflite_model_file = pathlib.Path('/tmp/model.tflite')  
2 tflite_model_file.write_bytes(tflite_model)
```

¹This is the standard for serializing a TensorFlow model.

1. Export the *SavedModel*¹:

```
1 export_dir = '/tmp/saved_model/1'  
2 tf.saved_model.save(model, export_dir=export_dir)
```

2. Convert the model:

```
1 converter = tf.lite.TFLiteConverter.from_saved_model(  
    export_dir)  
2 tflite_model = converter.convert()
```

3. Save the model:

```
1 tflite_model_file = pathlib.Path('/tmp/model.tflite')  
2 tflite_model_file.write_bytes(tflite_model)
```

¹This is the standard for serializing a TensorFlow model.

1. Export the *SavedModel*¹:

```
1 export_dir = '/tmp/saved_model/1'  
2 tf.saved_model.save(model, export_dir=export_dir)
```

2. Convert the model:

```
1 converter = tf.lite.TFLiteConverter.from_saved_model(  
    export_dir)  
2 tflite_model = converter.convert()
```

3. Save the model:

```
1 tflite_model_file = pathlib.Path('/tmp/model.tflite')  
2 tflite_model_file.write_bytes(tflite_model)
```

¹This is the standard for serializing a TensorFlow model.

1

Build a
model

2

Export and
convert

3

Verify

4

Deploy

TensorFlow Lite has a built-in interpreter in Python. You can test your model before deployment.

1. Load the TensorFlow Lite model and allocate tensors:

```
1 interpreter = tf.lite.Interpreter(model_path=model_path)
2 interpreter.allocate_tensors()
```

2. Get input and output tensors:

```
1 input_details = interpreter.get_input_details()
2 output_details = interpreter.get_output_details()
```

3. Test the model on input data:

```
1 interpreter.set_tensor(input_details[0]['index'],
    input_data)
```

TensorFlow Lite has a built-in interpreter in Python. You can test your model before deployment.

1. Load the TensorFlow Lite model and allocate tensors:

```
1 interpreter = tf.lite.Interpreter(model_path=model_path)
2 interpreter.allocate_tensors()
```

2. Get input and output tensors:

```
1 input_details = interpreter.get_input_details()
2 output_details = interpreter.get_output_details()
```

3. Test the model on input data:

```
1 interpreter.set_tensor(input_details[0]['index'],
    input_data)
```

TensorFlow Lite has a built-in interpreter in Python. You can test your model before deployment.

1. Load the TensorFlow Lite model and allocate tensors:

```
1 interpreter = tf.lite.Interpreter(model_path=model_path)
2 interpreter.allocate_tensors()
```

2. Get input and output tensors:

```
1 input_details = interpreter.get_input_details()
2 output_details = interpreter.get_output_details()
```

3. Test the model on input data:

```
1 interpreter.set_tensor(input_details[0]['index'],
    input_data)
```

4. Run the interpreter

```
1 interpreter.invoke()
```

5. Get the output data:

```
1 output_data = interpreter.get_tensor(output_details[0][ '
    index ' ])
```


4. Run the interpreter

```
1 interpreter.invoke()
```

5. Get the output data:

```
1 output_data = interpreter.get_tensor(output_details[0][ '
    index '])
```

DIY. Start with the notebook provided on .



☑ You can now use the model on your device!

1

Build a
model

2

Export and
convert

3

Verify

4

Deploy

Optimization

Optimization might be necessary **due to limited resources**.

With optimization, one wants to **reduce latency and the power consumption** of the models.

Solution ➕ Quantization!

Quantization

Is an optimization technique that **reduces the precision of the numbers** in the weights and biases of the model, but at the same time, **reduces model size and improves CPU and hardware accelerator latency**.

You will notice, however, a **little degradation in the model's accuracy**!

Optimization

Optimization might be necessary **due to limited resources**.

With optimization, one wants to **reduce latency and the power consumption** of the models.

Solution ➕ Quantization!

Quantization

Is an optimization technique that **reduces the precision of the numbers** in the weights and biases of the model, but at the same time, **reduces model size and improves CPU and hardware accelerator latency**.

You will notice, however, a **little degradation in the model's accuracy**!

Optimization

Optimization might be necessary due to limited resources.

With optimization, one wants to reduce latency and the power consumption of the models.

Solution ➕ Quantization!

Quantization

Is an optimization technique that reduces the precision of the numbers in the weights and biases of the model, but at the same time, reduces model size and improves CPU and hardware accelerator latency.

You will notice, however, a little degradation in the model's accuracy!

Optimization

Optimization might be necessary due to limited resources.

With optimization, one wants to reduce latency and the power consumption of the models.

Solution ➕ Quantization!

Quantization

Is an optimization technique that reduces the precision of the numbers in the weights and biases of the model, but at the same time, reduces model size and improves CPU and hardware accelerator latency.

You will notice, however, a little degradation in the model's accuracy!

Types of quantization

Dynamic

Quantizes **only the weights** from floating-point to integer, which has 8-bits of precision. However, inference operations are still done in floating-point precision.

Full integer quantization

This technique makes **all the model's math integer quantized**. However, it uses float operators when they don't have an integer implementation.

Integer only

Only used to **ensure compatibility with integer only devices** (such as 8-bit microcontrollers) by enforcing full-integer quantization for all operations including the input and output.

Types of quantization

Dynamic

Quantizes **only the weights** from floating-point to integer, which has 8-bits of precision. However, inference operations are still done in floating-point precision.

Full integer quantization

This technique makes **all the model's math integer quantized**. However, it uses float operators when they don't have an integer implementation.

Integer only

Only used to **ensure compatibility with integer only devices** (such as 8-bit microcontrollers) by enforcing full-integer quantization for all operations including the input and output.

Types of quantization

Dynamic

Quantizes **only the weights** from floating-point to integer, which has 8-bits of precision. However, inference operations are still done in floating-point precision.

Full integer quantization

This technique makes **all the model's math integer quantized**. However, it uses float operators when they don't have an integer implementation.

Integer only

Only used to **ensure compatibility with integer only devices** (such as 8-bit microcontrollers) by enforcing full-integer quantization for all operations including the input and output.

Quantization in code

You can specify the type of quantization technique that you want to use for your model in the converter options.

- **Dynamic**

```
1 converter.optimizations = [tf.lite.Optimize.DEFAULT]
```

- **Full integer quantization**

```
1 converter.optimizations = [tf.lite.Optimize.DEFAULT]
2 converter.representative_dataset = representative_dataset
```

- **Integer only**

```
1 converter.target_spec.supported_ops = [tf.lite.OpsSet.
    TFLITE_BUILTINS_INT8]
2 converter.inference_input_type = tf.int8
3 converter.inference_output_type = tf.int8
```

Quantization in code

You can specify the type of quantization technique that you want to use for your model in the converter options.

- **Dynamic**

```
1 converter.optimizations = [tf.lite.Optimize.DEFAULT]
```

- **Full integer quantization**

```
1 converter.optimizations = [tf.lite.Optimize.DEFAULT]
2 converter.representative_dataset = representative_dataset
```

- **Integer only**

```
1 converter.target_spec.supported_ops = [tf.lite.OpsSet.
    TFLITE_BUILTINS_INT8]
2 converter.inference_input_type = tf.int8
3 converter.inference_output_type = tf.int8
```

Quantization in code

You can specify the type of quantization technique that you want to use for your model in the converter options.

- **Dynamic**

```
1 converter.optimizations = [tf.lite.Optimize.DEFAULT]
```

- **Full integer quantization**

```
1 converter.optimizations = [tf.lite.Optimize.DEFAULT]
2 converter.representative_dataset = representative_dataset
```

- **Integer only**

```
1 converter.target_spec.supported_ops = [tf.lite.OpsSet.
    TFLITE_BUILTINS_INT8]
2 converter.inference_input_type = tf.int8
3 converter.inference_output_type = tf.int8
```

Quantization in code (cont.)

It's a small dataset to **calibrate or estimate the range**, i.e (min, max), of all floating-point arrays in the model for quantization.

This dataset **can be a small subset** (around 100–500 samples) of the **training or validation data**.

This dataset is generated using a **generator function** (remember the first day of the workshop?), such as:

```
1 def representative_dataset():
2     for _ in range(100):
3         data = x_train[:100]
4         yield [tf.dtypes.cast(data, tf.float32)]
```

It's a small dataset to **calibrate or estimate the range**, i.e (min, max), of all floating-point arrays in the model for quantization.

This dataset **can be a small subset** (around 100–500 samples) of the **training or validation data**.

This dataset is generated using a **generator function** (remember the first day of the workshop?), such as:

```
1 def representative_dataset():
2     for _ in range(100):
3         data = x_train[:100]
4         yield [tf.dtypes.cast(data, tf.float32)]
```

Quantization in code (cont.)


It's a small dataset to **calibrate or estimate the range**, i.e (min, max), of all floating-point arrays in the model for quantization.

This dataset **can be a small subset** (around 100–500 samples) of the **training or validation data**.

This dataset is generated using a **generator function** (remember the first day of the workshop?), such as:

```
1 def representative_dataset():
2     for _ in range(100):
3         data = x_train[:100]
4         yield [tf.dtypes.cast(data, tf.float32)]
```

DIY. Start with the notebook provided on .



TensorFlow Lite has still some limitations in terms of **operator compatibility**.

Some models are still **relatively too big** to store on devices.

One solution ➡ Execute the model in a remote server using a REST API!

REST API

TensorFlow Lite has still some limitations in terms of **operator compatibility**.

Some models are still **relatively too big** to store on devices.


One solution ➡ Execute the model in a remote server using a REST API!

REST API

TensorFlow Lite has still some limitations in terms of **operator compatibility**.

Some models are still **relatively too big** to store on devices.

One solution ➡ Execute the model in a remote server using a REST API!



- HTTP request by the client:

```
1 GET workshop/?input= and hell!\nAll hurt behind; backs red,  
   and faces pale\nWith f
```

- HTTP response by the remote server:

```
1 {  
2   "input": " and hell!\nAll hurt behind; backs red, and  
   faces pale\nWith f",  
3   "output": " and hell!\nAll hurt behind; backs red, and  
   faces pale\nWith fake wears theme your fore is man, no  
   'sbers; bays?\nGut that tho becues! Grown\nwet as a  
   teace paolly, hich.\n\nQUEEs I' meer theing that.\nHe  
   suatt mo carknsd suchenics, for Ratherrshads;\nIn  
   worvake furb\n"  
4 }
```

- HTTP request by the client:


```
1 GET workshop/?input= and hell!\nAll hurt behind; backs red ,  
   and faces pale\nWith f
```

- HTTP response by the remote server:

```
1 {  
2   "input": " and hell!\nAll hurt behind; backs red , and  
   faces pale\nWith f",  
3   "output": " and hell!\nAll hurt behind; backs red , and  
   faces pale\nWith fake wears theme your fore is man, no  
   'sbers; bays?\nGut that tho becues! Grown\nwet as a  
   teace paolly, hich.\n\nQUEEs I' meer theing that.\nHe  
   suatt mo carknsed suchenics, for Ratherrshads;\nIn  
   worvake furb\n"  
4 }
```

A **status code is issued** by the remote server in response to a client's request made to the server.

Code		Status
200		OK
400		Bad request
404		Not found
500	Internal server error	
⋮		⋮

DIY. Start with the notebook provided on .



Transfer Learning with TensorFlow Hub

TensorFlow Hub (<https://tfhub.dev>) is a **repository of trained machine learning models** ready to use.

These **models can be then fine-tuned for a specific problem** and deployed anywhere using, e.g., using TensorFlow Lite.

With transfer learning, a **model developed for a specific task can reused as the starting point (e.g., for feature extraction) for other model on a different task.**

Transfer Learning with TensorFlow Hub

TensorFlow Hub (<https://tfhub.dev>) is a repository of trained machine learning models ready to use.

These models can be then fine-tuned for a specific problem and deployed anywhere using, e.g., using TensorFlow Lite.


With transfer learning, a model developed for a specific task can be reused as the starting point (e.g., for feature extraction) for other model on a different task.


Transfer Learning with TensorFlow Hub

TensorFlow Hub (<https://tfhub.dev>) is a repository of trained machine learning models ready to use.

These models can be then fine-tuned for a specific problem and deployed anywhere using, e.g., using TensorFlow Lite.

With transfer learning, a model developed for a specific task can be reused as the starting point (e.g., for feature extraction) for other model on a different task.

DIY. Start with the notebook provided on .



Questions?